

Pattern Matching

Document #: P1371R0
Date: 2019-01-21
Project: Programming Language C++
Evolution
Reply-to: Sergei Murzin
<smurzin@bloomberg.net>
Michael Park
<mcypark@gmail.com>
David Sankel
<dsankel@bloomberg.net>
Dan Sarginson
<dsarginson@bloomberg.net>

Contents

1	Revision History	3
2	Introduction	3
3	Motivation and Scope	3
4	Before/After Comparisons	3
4.1	Matching Integrals	3
4.2	Matching Strings	4
4.3	Matching Tuples	4
4.4	Matching Variants	4
4.5	Matching Polymorphic Types	5
4.6	Evaluating Expression Trees	5
5	Design Overview	7
5.1	Basic Syntax	7
5.2	Basic Model	7
5.3	Types of Patterns	8
5.3.1	Primary Patterns	8
5.3.1.1	Wildcard Pattern	8
5.3.1.2	Identifier Pattern	8
5.3.1.3	Expression Pattern	8
5.3.2	Compound Patterns	9
5.3.2.1	Structured Binding Pattern	9
5.3.2.2	Alternative Pattern	10
5.3.2.3	Parenthesized Pattern	13
5.3.2.4	Binding Pattern	13
5.3.2.5	Dereference Pattern	13
5.3.2.6	Extractor Pattern	14
5.4	Pattern Guard	15

5.5	<code>inspect constexpr</code>	15
5.6	Exhaustiveness and Usefulness	16
5.7	Refutability	16
6	Proposed Wording	16
7	Design Decisions	17
7.1	Extending Structured Bindings Declaration	17
7.2	<code>inspect</code> rather than <code>switch</code>	17
7.3	First Match rather than Best Match	17
7.4	Unrestricted Side Effects	18
7.5	Language rather than Library	19
7.6	Matchers and Extractors	19
7.7	Wildcard Syntax	19
7.8	Expression vs Pattern Disambiguation	19
8	Runtime Performance	21
8.1	Structured Binding Pattern	21
8.2	Alternative Pattern	21
8.3	Open Class Hierarchy	22
9	Examples	22
9.1	Predicate-based Discriminator	22
9.2	“Closed” Class Hierarchy	23
9.3	Matcher: <code>any_of</code>	24
9.4	Matcher: <code>within</code>	25
9.5	Extractor: <code>both</code>	25
9.6	Extractor: <code>at</code>	25
9.7	Red-black Tree Rebalancing	26
10	Future Work	28
10.1	Language Support for Variant	28
10.2	Note on Ranges	28
11	Acknowledgements	28
12	References	29

1 Revision History

- R0 — Merged [\[P1260R0\]](#) and [\[P1308R0\]](#)

2 Introduction

As algebraic data types gain better support in C++ with facilities such as `tuple` and `variant`, the importance of mechanisms to interact with them have increased. While mechanisms such as `apply` and `visit` have been added, their usage is quite complex and limited even for simple cases. Pattern matching is a widely adopted mechanism across many programming languages to interact with algebraic data types that can help greatly simplify C++. Examples of programming languages include text-based languages such as SNOBOL back in the 1960s, functional languages such as Haskell and OCaml, and “mainstream” languages such as Scala, Swift, and Rust.

This paper is a result of collaboration between the authors of [\[P1260R0\]](#) and [\[P1308R0\]](#). A joint presentation by the authors of the two proposals was given in EWGI at the San Diego 2018 meeting, with the closing poll: “Should we commit additional committee time to pattern matching?” — SF: 14, WF: 0, N: 1, WA: 0, SA: 0

3 Motivation and Scope

Virtually every program involves branching on some predicates applied to a value and conditionally binding names to some of its components for use in subsequent logic. Today, C++ provides two types of selection statements: the `if` statement and the `switch` statement.

Since `switch` statements can only operate on a *single* integral value and `if` statements operate on an *arbitrarily* complex boolean expression, there is a significant gap between the two constructs even in inspection of the “vocabulary types” provided by the standard library.

In C++17, structured binding declarations [\[P0144R2\]](#) introduced the ability to concisely bind names to components of `tuple`-like values. The proposed direction of this paper aims to naturally extend this notion by performing **structured inspection** with `inspect` statements and expressions. The goal of `inspect` is to bridge the gap between `switch` and `if` statements with a **declarative**, **structured**, **cohesive**, and **composable** mechanism.

4 Before/After Comparisons

4.1 Matching Integrals

Before

```
switch (x) {
  case 0: std::cout << "got zero"; break;
  case 1: std::cout << "got one"; break;
  default: std::cout << "don't care";
}
```

After

```
inspect (x) {
  0: std::cout << "got zero";
  1: std::cout << "got one";
  _: std::cout << "don't care";
}
```

4.2 Matching Strings

Before

```
if (s == "foo") {
    std::cout << "got foo";
} else if (s == "bar") {
    std::cout << "got bar";
} else {
    std::cout << "don't care";
}
```

After

```
inspect (s) {
    "foo": std::cout << "got foo";
    "bar": std::cout << "got bar";
    _: std::cout << "don't care";
}
```

4.3 Matching Tuples

Before

```
auto&& [x, y] = p;
if (x == 0 && y == 0) {
    std::cout << "on origin";
} else if (x == 0) {
    std::cout << "on y-axis";
} else if (y == 0) {
    std::cout << "on x-axis";
} else {
    std::cout << x << ',' << y;
}
```

After

```
inspect (p) {
    [0, 0]: std::cout << "on origin";
    [0, y]: std::cout << "on y-axis";
    [x, 0]: std::cout << "on x-axis";
    [x, y]: std::cout << x << ',' << y;
}
```

4.4 Matching Variants

Before

```
struct visitor {
    void operator()(int i) const {
        os << "got int: " << i;
    }
    void operator()(float f) const {
        os << "got float: " << f;
    }
    std::ostream& os;
};
std::visit(visitor{strm}, v);
```

After

```
inspect (v) {
    <int> i: strm << "got int: " << i;
    <float> f: strm << "got float: " << f;
}
```

4.5 Matching Polymorphic Types

```
struct Shape { virtual ~Shape() = default; };
struct Circle : Shape { int radius; };
struct Rectangle : Shape { int width, height; };
```

Before

```
virtual int Shape::get_area() const = 0;

int Circle::get_area() const override {
    return 3.14 * radius * radius;
}

int Rectangle::get_area() const override {
    return width * height;
}
```

After

```
int get_area(const Shape& shape) {
    return inspect (shape) {
        <Circle>    [r]    => 3.14 * r * r,
        <Rectangle> [w, h] => w * h
    }
}
```

4.6 Evaluating Expression Trees

```
struct Expr;

struct Neg {
    std::shared_ptr<Expr> expr;
};

struct Add {
    std::shared_ptr<Expr> lhs, rhs;
};

struct Mul {
    std::shared_ptr<Expr> lhs, rhs;
};

struct Expr : std::variant<int, Neg, Add, Mul> {
    using variant::variant;
};

namespace std {
    template <>
    struct variant_size<Expr> : variant_size<Expr::variant> {};

    template <std::size_t I>
    struct variant_alternative<I, Expr> : variant_alternative<I, Expr::variant> {};
}
```

Before

```
int eval(const Expr& expr) {
    struct visitor {
        int operator()(int i) const {
            return i;
        }
        int operator()(const Neg& n) const {
            return -eval(*n.expr);
        }
        int operator()(const Add& a) const {
            return eval(*a.lhs) + eval(*a.rhs);
        }
        int operator()(const Mul& m) const {
            // Optimize multiplication by 0.
            if (int* i = std::get_if<int>(m.lhs.get()); i && *i == 0) {
                return 0;
            }
            if (int* i = std::get_if<int>(m.rhs.get()); i && *i == 0) {
                return 0;
            }
            return eval(*m.lhs) * eval(*m.rhs);
        }
    };
    return std::visit(visitor{}, expr);
}
```

After

```
int eval(const Expr& expr) {
    inspect (expr) {
        <int> i: return i;
        <Neg> [*e]: return -eval(e);
        <Add> [*l, *r]: return eval(l) + eval(r);
        // Optimize multiplication by 0.
        <Mul> [*(<int> 0), _]: return 0;
        <Mul> [_, *(<int> 0)]: return 0;
        <Mul> [*l, *r]: return eval(l) * eval(r);
    }
}
```

5 Design Overview

5.1 Basic Syntax

There are two forms of `inspect`: the statement form and the expression form.

```
inspect constexpropt ( init-statementopt condition ) {  
    pattern guardopt : statement  
    pattern guardopt : statement  
    ...  
}  
  
inspect constexpropt ( init-statementopt condition ) trailing-return-typeopt {  
    pattern guardopt => expression ,  
    pattern guardopt => expression ,  
    ...  
}  
  
guard:  
    if ( expression )
```

[*Note*: The expression form is roughly equivalent to:

```
std::invoke([&]() trailing-return-typeopt {  
    inspect constexpropt ( init-statementopt condition ) {  
        pattern guardopt : return expression ;  
        pattern guardopt : return expression ;  
        ...  
    }  
})
```

— *end note*]

5.2 Basic Model

Within the parentheses, the `inspect` statement is equivalent to `switch` and `if` statements except that no conversion nor promotion takes place in evaluating the value of its condition.

When the `inspect` statement is executed, its condition is evaluated and matched in order (first match semantics) against each pattern. If a pattern successfully matches the value of the condition and the boolean expression in the guard evaluates to `true` (or if there is no guard at all), control is passed to the statement following the matched pattern label. If the guard expression evaluates to `false`, control flows to the subsequent pattern.

If no pattern matches, none of the statements are executed for the statement form and `std::no_match` exception is thrown for the expression form.

5.3 Types of Patterns

5.3.1 Primary Patterns

5.3.1.1 Wildcard Pattern

The wildcard pattern has the form:

```
-  
and matches any value v.  
int v = /* ... */;  
  
inspect (v) {  
    _: std::cout << "ignored";  
    // ^ wildcard pattern  
}
```

[*Note*: Refer to [Wildcard Syntax](#) for a design discussion. — *end note*]

5.3.1.2 Identifier Pattern

The identifier pattern has the form:

identifier

and matches any value *v*. The introduced name behaves as an lvalue referring to *v*, and is in scope from its point of declaration until the end of the statement following the pattern label.

```
int v = /* ... */;  
  
inspect (v) {  
    x: std::cout << x;  
    // ^ identifier pattern  
}
```

[*Note*: If the identifier pattern is used at the top-level, it has the same syntax as a `goto` label. — *end note*]

5.3.1.3 Expression Pattern

The expression pattern has the form:

literal
this
^ primary-expression

and matches value *v* if a call to member `e.match(v)` or else a non-member ADL-only `match(e, v)` is contextually convertible to `bool` and evaluates to `true` where *e* is the *literal* or *primary-expression*.

The default behavior of `match(x, y)` is `x == y`.


```

int v = /* ... */;

inspect (v) {
    0: std::cout << "got zero";
    1: std::cout << "got one";
    // ^ expression pattern
}

static constexpr int zero = 0, one = 1;
int v = /* ... */;

inspect (v) {
    ^zero: std::cout << "got zero";
    // ^^^^^ expression pattern
}

```

5.3.2 Compound Patterns

5.3.2.1 Structured Binding Pattern

The structured binding pattern has the following two forms:

```

[ pattern0 , pattern1 , ... , patternN ]
[ designator0 : pattern0 , designator1 : pattern1 , ... , designatorN : patternN ]

```

The first form matches value *v* if each *pattern*_{*i*} matches the *i*th component of *v*. The components of *v* are given by the structured binding declaration: `auto&& [__e0, __e1, ..., __eN] = v;` where each __e_{*i*} are unique exposition-only identifiers.

```

std::pair<int, int> p = /* ... */;

inspect (p) {
    [0, 0]: std::cout << "on origin";
    [0, y]: std::cout << "on y-axis";
    // ^ identifier pattern
    [x, 0]: std::cout << "on x-axis";
    // ^ expression pattern
    [x, y]: std::cout << x << ', ' << y;
    // ^^^^^ structured binding pattern
}

```

The second form matches value *v* if each *pattern*_{*i*} matches the direct non-static data member of *v* named *identifier* from each *designator*_{*i*}. If an *identifier* from any *designator*_{*i*} does not refer to a direct non-static data member of *v*, the program is ill-formed.

```

struct Player { std::string name; int hitpoints; int coins; };

void get_hint(const Player& p) {
    inspect (p) {
        [.hitpoints: 1]: std::cout << "You're almost destroyed. Give up!\n";
        [.hitpoints: 10, .coins: 10]: std::cout << "I need the hints from you!\n";
        [.coins: 10]: std::cout << "Get more hitpoints!\n";
        [.hitpoints: 10]: std::cout << "Get more ammo!\n";
        [.name: n]: {
            if (n != "The Bruce Dickenson") {
                std::cout << "Get more hitpoints and ammo!\n";
            } else {
                std::cout << "More cowbell!\n";
            }
        }
    }
}

```

[*Note*: Unlike designated initializers, the order of the designators need not be the same as the declaration order of the members of the class. — *end note*]

5.3.2.2 Alternative Pattern

The alternative pattern has the following forms:

```

< auto > pattern
< concept > pattern
< type > pattern
< constant-expression > pattern

```

Let *v* be the value being matched and *V* be `std::remove_cvref_t<decltype(v)>`.

Let *Alt* be the entity inside the angle brackets.

Case 1: `std::variant-like`

If `std::variant_size_v<V>` is well-formed and evaluates to an integral, the alternative pattern matches *v* if *Alt* is compatible with the current index of *v* and *pattern* matches the active alternative of *v*.

Let *I* be the current index of *v* given by a member *v.index()* or else a non-member ADL-only *index(v)*. The active alternative of *v* is given by `std::variant_alternative_t<I, V>&` initialized by a member *v.get<I>()* or else a non-member ADL-only *get<I>(v)*.

Alt is compatible with *I* if one of the following four cases is true:

- *Alt* is `auto`
- *Alt* is a *concept* and `std::variant_alternative_t<I, V>` satisfies the *concept*.
- *Alt* is a *type* and `std::is_same_v<Alt, std::variant_alternative_t<I, V>>` is `true`
- *Alt* is a *constant-expression* that can be used in a `switch` and is the same value as *I*.

Before

```

std::visit([&](auto&& x) {
    strm << "got auto: " << x;
}, v);

```

After

```

inspect (v) {
    <auto> x: strm << "got auto: " << x;
}

```

```
std::visit([&](auto&& x) {
    using X = std::remove_cvref_t<decltype(x)>;
    if constexpr (C1<X>()) {
        strm << "got C1: " << x;
    } else if constexpr (C2<X>()) {
        strm << "got C2: " << x;
    }
}, v);
```

```
inspect (v) {
    <C1> c1: strm << "got C1: " << c1;
    <C2> c2: strm << "got C2: " << c2;
}
```

```
std::visit([&](auto&& x) {
    using X = std::remove_cvref_t<decltype(x)>;
    if constexpr (std::is_same_v<int, X>) {
        strm << "got int: " << x;
    } else if constexpr (
        std::is_same_v<float, X>) {
        strm << "got float: " << x;
    }
}, v);
```

```
inspect (v) {
    <int> i: strm << "got int: " << i;
    <float> f: strm << "got float: " << f;
}
```

```
std::variant<int, int> v = /* ... */;

std::visit([&](int x) {
    strm << "got int: " << x;
}, v);
```

```
std::variant<int, int> v = /* ... */;

inspect (v) {
    <int> x: strm << "got int: " << x;
}
```

```
std::variant<int, int> v = /* ... */;

std::visit([&](auto&& x) {
    switch (v.index()) {
        case 0: {
            strm << "got first: " << x;
            break;
        }
        case 1: {
            strm << "got second: " << x;
            break;
        }
    }
}, v);
```

```
std::variant<int, int> v = /* ... */;

inspect (v) {
    <0> x: strm << "got first: " << x;
    <1> x: strm << "got second: " << x;
}
```

Case 2: std::any-like

< type > pattern

If *Alt* is a *type* and there exists a valid non-member ADL-only `any_cast<Alt>(&v)`, let *p* be its result. The alternative pattern matches if *p* contextually converted to `bool` evaluates to `true`, and *pattern* matches **p*.

Before

```
std::any a = 42;

if (int* i = any_cast<int>(&a)) {
    std::cout << "got int: " << *i;
} else if (float* f = any_cast<float>(&a)) {
    std::cout << "got float: " << *f;
}
```

After

```
std::any a = 42;

inspect (a) {
    <int> i: std::cout << "got int: " << i;
    <float> f: std::cout << "got float: " << f;
}
```

Case 3: Polymorphic Types

< type > pattern

If *Alt* is a *type* and `std::is_polymorphic_v<V>` is `true`, let *p* be `dynamic_cast<Alt'*>(&v)` where *Alt'* has the same *cv*-qualifications as `decltype(&v)`. The alternative pattern matches if *p* contextually converted to `bool` evaluates to `true`, and *pattern* matches **p*.

While the **semantics** of the pattern is specified in terms of `dynamic_cast`, [N3449] describes techniques involving vtable pointer caching and hash conflict minimization that are implemented in the [Mach7] library, as well as mentions of further opportunities available for a compiler intrinsic.

Given the following definition of a `Shape` class hierarchy:

```
struct Shape { virtual ~Shape() = default; };

struct Circle : Shape { int radius; };
struct Rectangle : Shape { int width, height; };
```

Before

```
virtual int Shape::get_area() const = 0;

int Circle::get_area() const override {
    return 3.14 * radius * radius;
}

int Rectangle::get_area() const override {
    return width * height;
}
```

After

```
int get_area(const Shape& shape) {
    inspect (shape) {
        <Circle> [r]: return 3.14 * r * r;
        <Rectangle> [w, h]: return w * h;
    }
}
```

5.3.2.3 Parenthesized Pattern

The parenthesized pattern has the form:

(pattern)

and matches value *v* if *pattern* matches it.

```
std::variant<Point, /* ... */> v = /* ... */;

inspect (v) {
    <Point> ([x, y]): // ...
    //      ~~~~~ parenthesized pattern
}
```

5.3.2.4 Binding Pattern

The binding pattern has the form:

identifier @ pattern

and matches value *v* if *pattern* matches it. The introduced name behaves as an lvalue referring to *v*, and is in scope from its point of declaration until the end of the statement following the pattern label.

```
std::variant<Point, /* ... */> v = /* ... */;

inspect (v) {
    <Point> (p @ [x, y]): // ...
    //      ~~~~~ binding pattern
}
```

5.3.2.5 Dereference Pattern

The dereference pattern has the form:

** pattern*

and matches value *v* if *v* is contextually convertible to `bool` and evaluates to `true`, and *pattern* matches `*v`.

```
struct Node {
    int value;
    std::unique_ptr<Node> lhs, rhs;
};

template <typename Visitor>
void print_leftmost(const Node& node) {
    inspect (node) {
        [.value: v, .lhs: nullptr]: std::cout << v << '\n';
        [.lhs: *l]: print_leftmost(l);
    //      ^^ dereference pattern
    }
}
```

[*Note:* Refer to [Red-black Tree Rebalancing](#) for a more complex example. — *end note*]

5.3.2.6 Extractor Pattern

The extractor pattern has the following two forms:

```
( constant-expression ! pattern )  
( constant-expression ? pattern )
```

Let *c* be the *constant-expression*. The first form matches value *v* if *pattern* matches *e* where *e* is the result of a call to member *c.extract(v)* or else a non-member ADL-only *extract(c, v)*.

```
template <typename T>  
struct Is {  
    template <typename Arg>  
    Arg&& extract(Arg&& arg) const {  
        static_assert(std::is_same_v<T, std::remove_cvref_t<Arg>>);  
        return std::forward<Arg>(arg);  
    }  
};  
  
template <typename T>  
inline constexpr Is<T> is;  
  
// P0480: `auto&& [std::string s, int i] = f();`  
inspect (f()) {  
    [(is<std::string>! s), (is<int>! i)]: // ...  
    // ~~~~~ extractor pattern  
}
```

For second form, let *e* be the result of a call to member *c.try_extract(v)* or else a non-member ADL-only *try_extract(c, v)*. It matches value *v* if *e* is contextually convertible to *bool*, evaluates to *true*, and *pattern* matches **e*.

```
struct Email {  
    std::optional<std::array<std::string_view, 2>>  
    try_extract(std::string_view sv) const;  
};  
  
inline constexpr Email email;  
  
struct PhoneNumber {  
    std::optional<std::array<std::string_view, 3>>  
    try_extract(std::string_view sv) const;  
};  
  
inline constexpr PhoneNumber phone_number;  
  
inspect (s) {  
    (email? [address, domain]): std::cout << "got an email";  
    (phone_number? ["415", _, _]): std::cout << "got a San Francisco phone number";  
    // ~~~~~ extractor pattern  
}
```

5.4 Pattern Guard

The pattern guard has the form:

```
if ( expression )
```

Let *e* be the result of *expression* contextually converted to `bool`. If *e* is `true`, control is passed to the corresponding statement. Otherwise, control flows to the subsequent pattern.

The pattern guard allows to perform complex tests that cannot be performed within the *pattern*. For example, performing tests across multiple bindings:

```
inspect (p) {  
  [x, y] if (test(x, y)): std::cout << x << ',' << y << " passed";  
  //      ~~~~~ pattern guard  
}
```

This also diminishes the desire for fall-through semantics within the statements, an unpopular feature even in `switch` statements.

5.5 inspect constexpr

Every *pattern* is able to determine whether it matches value *v* as a boolean expression in isolation. Let `MATCHES` be the condition for which a *pattern* matches a value *v*. Ignoring any potential optimization opportunities, we're able to perform the following transformation:

inspect

```
inspect (v) {  
  pattern1 if (cond1): stmt1  
  pattern2: stmt2  
  // ...  
}
```

if

```
if (MATCHES(pattern1, v) && cond1) stmt1  
else if (MATCHES(pattern2, v)) stmt2  
// ...
```

`inspect constexpr` is then formulated by applying `constexpr` to every `if` branch.

inspect constexpr

```
inspect constexpr (v) {  
  pattern1 if (cond1): stmt1  
  pattern2: stmt2  
  // ...  
}
```

if constexpr

```
if constexpr (MATCHES(pattern1, v) && cond1) stmt1  
else if constexpr (MATCHES(pattern2, v)) stmt2  
// ...
```

5.6 Exhaustiveness and Usefulness

`inspect` can be declared `[[strict]]` for implementation-defined exhaustiveness and usefulness checking.

Exhaustiveness means that all values of the type of the value being matched is handled by at least one of the cases. For example, having a `_:` case makes any `inspect` statement exhaustive.

Usefulness means that every case handles at least one value of the type of the value being matched. For example, any case that comes after a `_:` case would be useless.

Warnings for pattern matching [\[Warnings\]](#) discusses and outlines an algorithm for exhaustiveness and usefulness for OCaml, and is the algorithm used by Rust.

5.7 Refutability

Patterns that cannot fail to match are said to be *irrefutable* in contrast to *refutable* patterns which can fail to match. For example, the identifier pattern is *irrefutable* whereas the expression pattern is *refutable*.

The distinction is useful in reasoning about which patterns should be allowed in which contexts. For example, the structured bindings declaration is conceptually a restricted form of pattern matching. With the introduction of expression pattern in this paper, some may question whether structured bindings declaration should be extended for examples such as `auto [0, x] = f();`.

This is ultimately a question of whether structured bindings declaration supports *refutable* patterns or if it is restricted to *irrefutable* patterns.

6 Proposed Wording

The following is the beginning of an attempt at a syntactic structure.

Add to §8.4 `[stmt.select]` of ...

- ¹ Selection statements choose one of several flows of control.

selection-statement:

```
if constexpropt ( init-statementopt condition ) statement
if constexpropt ( init-statementopt condition ) statement else statement
switch ( init-statementopt condition ) statement
inspect constexpropt ( init-statementopt condition ) trailing-return-typeopt { inspect-case-seq }
```

inspect-case-seq:

```
inspect-statement-case-seq
inspect-expression-case-seq
```

inspect-statement-case-seq:

```
inspect-statement-case
inspect-statement-case-seq inspect-statement-case
```

inspect-expression-case-seq:

```
inspect-expression-case
inspect-expression-case-seq , inspect-expression-case
```

inspect-statement-case:

```
inspect-pattern inspect-guardopt : statement
```



```

inspect-expression-case:
  inspect-pattern inspect-guardopt => assignment-expression

inspect-pattern:
  wildcard-pattern
  identifier-pattern
  expression-pattern
  structured-binding-pattern
  alternative-pattern
  binding-pattern
  dereference-pattern
  extractor-pattern

inspect-guard:
  if ( expression )

```

7 Design Decisions

7.1 Extending Structured Bindings Declaration

The design is intended to be consistent and to naturally extend the notions introduced by structured bindings. That is, The subobjects are **referred** to rather than being assigned into new variables.

7.2 `inspect` rather than `switch`

This proposal introduces a new `inspect` statement rather than trying to extend the `switch` statement. [P0095R0] had proposed extending `switch` and received feedback to “leave `switch` alone” in Kona 2015.

The following are some of the reasons considered:

- `switch` allows the `case` labels to appear **anywhere**, which hinders the goal of pattern matching in providing **structured** inspection.
- The fall-through semantics of `switch` generally results in `break` being attached to every case, and is known to be error-prone.
- `switch` is purposely restricted to integrals for **guaranteed** efficiency. The primary goal of pattern matching in this paper is expressiveness while being at least as efficient as the naively hand-written code.

7.3 First Match rather than Best Match

The proposed matching algorithm has first match semantics. The choice of first match is mainly due to complexity. Our overload resolution rules for function declarations are extremely complex and is often a mystery.

Best match via overload resolution for function declarations are absolutely necessary due to the non-local and unordered nature of declarations. That is, function declarations live in different files and get pulled in via mechanisms such as `#include` and `using` declarations, and there is no defined order of declarations like Haskell does, for example. If function dispatching depended on the order of `#include` and/or `using` declarations being pulled in from hundreds of files, it would be a complete disaster.

Pattern matching on the other hand do not have this problem because the construct is local and ordered in nature. That is, all of the candidate patterns appear locally within `inspect (x) { /* ... */ }` which cannot span across multiple files, and appear in a specified order. This is consistent with `try/catch` for the same reasons: locality and order.

Consider also the amount of limitations we face in overload resolution due to the opacity of user-defined types. `T*` is related to `unique_ptr<T>` as it is to `vector<T>` as far as the type system is concerned. This limitation will likely be even bigger in a pattern matching context with the amount of customization points available for user-defined behavior.

7.4 Unrestricted Side Effects

We considered the possibility of restricting side-effects within patterns. Specifically whether modifying the value currently being matched in the middle of evaluation should have defined behavior.

The consideration was due to potential optimization opportunities.

```
void f(int &); // defined in a different translation unit.
int x = 1;

inspect (x) {
  0: std::cout << 0;
  1 if (f(x)): std::cout << 1;
  2: std::cout << 2;
}
```

If modifying the value currently being matched has undefined behavior, a compiler can assume that `f` (defined in a different translation unit) will not change the value of `x`. This means that the compiler can generate code that uses a jump table to determine which of the patterns match.

If on the other hand `f` may change the value of `x`, the compiler would be forced to generated code checks the patterns in sequence, since a subsequent pattern may match the updated value of `x`.

The following are **illustrations** of the two approaches written in C++:

Not allowed to modify	Allowed to modify
<pre>void f(int &); int x = 1; switch (x) { case 0: std::cout << 0; break; case 1: if (f(x)) { std::cout << 1; } break; case 2: std::cout << 2; break; }</pre>	<pre>void f(int &); int x = 1; if (x == 0) std::cout << 0; else if (x == 1 && f(x)) std::cout << 1; else if (x == 2) std::cout << 2;</pre>

However, we consider this opportunity too niche. Suppose we have a slightly more complex case: `struct S { int x; }; and bool operator==(const S&, const S&);. Even if modifying the value being matched has undefined behavior, if the operator== is defined in a different translation unit, a compiler cannot do much more than generate code that checks the patterns in sequence anyway.`

7.5 Language rather than Library

There are three popular pattern matching libraries for C++ today: [\[Mach7\]](#), [\[Patterns\]](#), and [\[SimpleMatch\]](#).

While the libraries have been useful for gaining experience with interfaces and implementation, the issue of introducing identifiers, syntactic overhead of the patterns, and the reduced optimization opportunities justify support as a language feature from a usability standpoint.

7.6 Matchers and Extractors

Many languages provide a wide array of patterns through various syntactic forms. While this is a potential direction for C++, it would mean that every new type of matching requires new syntax to be added to the language. This would result in a narrow set of types being supported through limited customization points.

Matchers and extractors are supported in order to minimize the number of patterns with special syntax. The following are example matchers and extractors that commonly have special syntax in other languages.

Matchers / Extractors	Other Languages
<code>any_of{1, 2, 3}</code>	<code>1 2 3</code>
<code>within{1, 10}</code>	<code>1..10</code>
<code>(both! [[x, 0], [0, y]])</code>	<code>[x, 0] & [0, y]</code>
<code>(at! [p, [x, y]])</code>	<code>p @ [x, y]</code>

Each of the matchers and extractors can be found in the [Examples](#) section.

7.7 Wildcard Syntax

[\[P1110R0\]](#) discusses wildcard/placeholder syntax in different contexts of the language. The intent of this proposal is to consolidate with the results of the decision of [\[P1110R0\]](#).

Even though `_` is a valid identifier, it does not introduce a name as doing so would result in redeclaration errors in the case where multiple wildcard `_` identifiers are used.

It is possible for users to have already introduced `_` as a type or variable name in the same scope where an `inspect` statement is used. As a highly-visible example, the authors are aware of the use of `_` as a name in the popular “Google Mock” library. Idiomatically this is accessed by introducing `_` into the current namespace or block scope with a `using` declaration. Using the wildcard pattern in cases like this is unambiguous since the expression pattern requires a `^` introducer for primary expressions. `^_` will always match against an existing name and `_` will always represent the wildcard pattern. An existing `_` name can be used without ambiguity in the matched statement to which control is passed.

Naturally, the impact of defining `_` in the pre-processor cannot be predicted or controlled by this paper and is thus liable to result in an ill-formed program.

7.8 Expression vs Pattern Disambiguation

Crucial discussion revolves around whether arbitrary expressions should be allowed inside patterns. Allowing expressions would provide the ability to compute values to match without creating variables for them before `inspect`. However, this introduces syntactic ambiguity problems that must be addressed.

Expressions can contain a vast number of syntactic elements, such as arithmetic, bit-shift, and logical operators, parentheses and braces, etc. This means that if we allow expressions to appear as patterns without disambiguation, we would be closing a significant amount of the syntax options for patterns. For example, consider the logical-or pattern that matches if any of given patterns match. Most languages spell it as $pattern_0 \mid pattern_1 \mid \dots \mid pattern_N$. However, examples such as $1 \mid 2$ becomes ambiguous as to whether $1 \mid 2$ is an expression (3) or a pattern (matches 1 or 2).

This paper proposes to allow arbitrary expression to appear in patterns (See [Expression Pattern](#)). We propose to use $\hat{}$ symbol to disambiguate expressions. This allows full flexibility of expressions to be used inline, while keeping the syntax space open for pattern development.

For example, let's look at identifier and expression patterns. Both of them can contain an *unqualified-id*. It can perform unqualified name lookup in the case of [Expression Pattern](#) and introduce a new identifier into scope in case of [Identifier Pattern](#).

```
constexpr int x = /* ... */

inspect (v) {
    ^x: std::cout << "`v` matched `x` by value";
    x: std::cout << "introduces new identifier `x` bound to `v`: " << x;
}
```

\hat{x} performs unqualified name lookup for x , while x introduces new identifier x in scope.

The problem described spans more than just *identifier pattern*. That could have been solved with an introducer for *identifier pattern* only. For example, [Swift Patterns](#) use keyword `let` to denote identifier pattern and treat everything else as an expression. This solution works but closes off pattern syntax that already looks like an expression. Using $1 \mid 2$ as an example again, it could not be introduced as a logical-or pattern in Swift without a breaking change since it's already a valid expression.

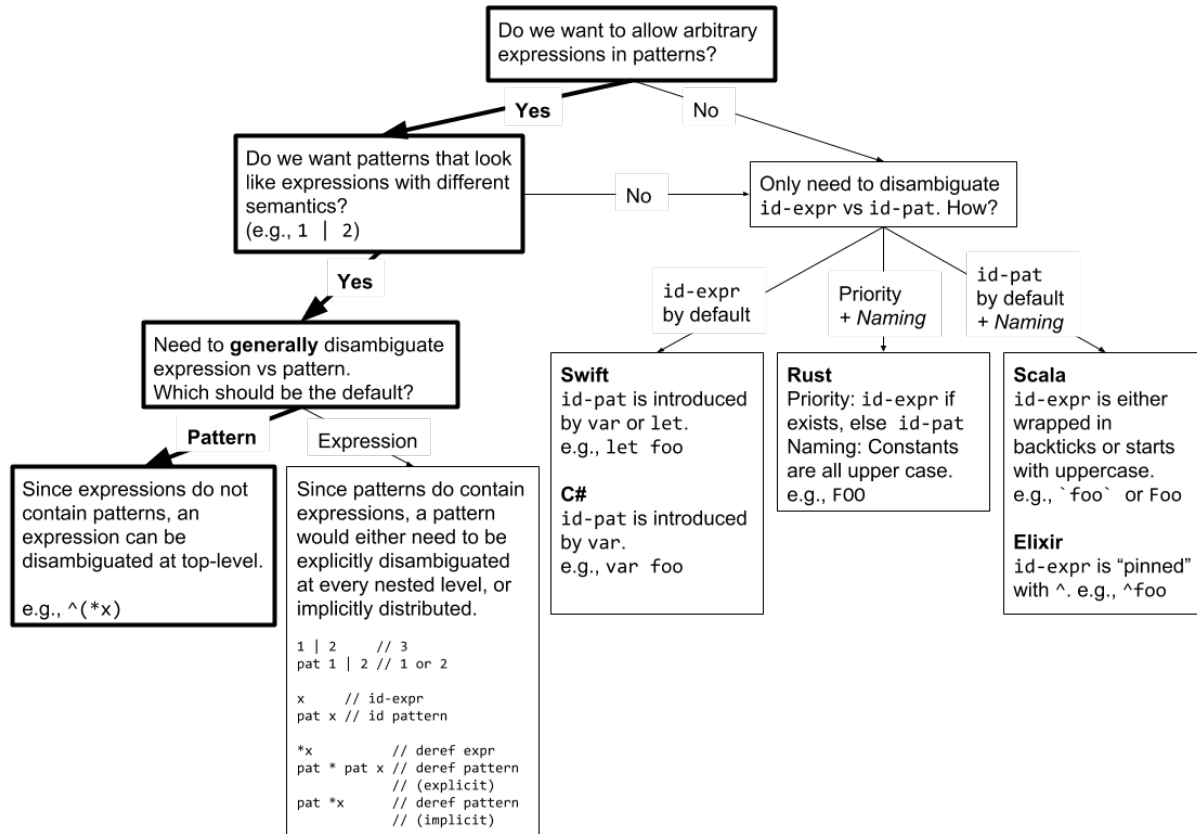
That leaves the only other solution — to disambiguate patterns rather than expressions. For example, we can prefix patterns with a context-sensitive keyword `pat`. Using the example from [Matching Tuples](#):

```
std::pair<int, int> p = /* ... */

inspect (p) {
    pat [0, 0]: std::cout << "on origin";
    pat [0, pat y]: std::cout << "on y-axis";
    pat [pat x, 0]: std::cout << "on x-axis";
    pat [pat x, pat y]: std::cout << x << ',' << y;
}
```

This solution however is more complex because patterns contain expressions. This means that a pattern would either need to be explicitly prefixed at nested level, or implicitly distributed to inner patterns.

The following is a flow graph of decisions that need to be made:



8 Runtime Performance

The following are few of the optimizations that are worth noting.

8.1 Structured Binding Pattern

Structured binding patterns can be optimized by performing `switch` over the columns with the duplicates removed, rather than the naive approach of performing a comparison per element. This removes unnecessary duplicate comparisons that would be performed otherwise. This would likely require some wording around “comparison elision” in order to enable such optimizations.

8.2 Alternative Pattern

The sequence of alternative patterns can be executed in a `switch`.

8.3 Open Class Hierarchy

[N3449] describes techniques involving vtable pointer caching and hash conflict minimization that are implemented in the [Mach7] library, but also mentions further opportunities available for a compiler solution.

9 Examples

9.1 Predicate-based Discriminator

Short-string optimization using a **predicate** as a discriminator rather than an explicitly stored **value**. Adapted from Bjarne Stroustrup's pattern matching presentation at Urbana-Champaign 2014 [PatMatPres].

```
struct String {
    enum Storage { Local, Remote };

    int size;
    union {
        char local[32];
        struct { char *ptr; int unused_allocated_space; } remote;
    };

    // Predicate-based discriminator derived from `size`.
    Storage index() const { return size > sizeof(local) ? Remote : Local; }

    // Opt into Variant-Like protocol.
    template <Storage S>
    auto &&get() {
        if constexpr (S == Local) return local;
        else if constexpr (S == Remote) return remote;
    }

    char *data();
};

namespace std {
    // Opt into Variant-Like protocol.

    template <>
    struct variant_size<String> : std::integral_constant<std::size_t, 2> {};

    template <>
    struct variant_alternative<String::Local, String> {
        using type = decltype(String::local);
    };

    template <>
    struct variant_alternative<String::Remote, String> {
        using type = decltype(String::remote);
    };
}
```

```

char* String::data() {
    inspect (*this) {
        <Local> l: return l;
        <Remote> r: return r.ptr;
    }
    // switch (index()) {
    //     case Local: {
    //         std::variant_alternative_t<Local, String>& l = get<Local>();
    //         return l;
    //     }
    //     case Remote: {
    //         std::variant_alternative_t<Remote, String>& r = get<Remote>();
    //         return r.ptr;
    //     }
    // }
}

```

9.2 “Closed” Class Hierarchy

A class hierarchy can effectively be closed with an `enum` that maintains the list of its members, and provide efficient dispatching by opting into the Variant-Like protocol.

A generalized mechanism of pattern is used extensively in LLVM; `llvm/Support/YAMLParse.h` [\[YAMLParse\]](#) is an example.

```

struct Shape { enum Kind { Circle, Rectangle } kind; };

struct Circle : Shape {
    Circle(int radius) : Shape{Shape::Kind::Circle}, radius(radius) {}

    int radius;
};

struct Rectangle : Shape {
    Rectangle(int width, int height)
        : Shape{Shape::Kind::Rectangle}, width(width), height(height) {}

    int width, height;
};

namespace std {
    template <>
    struct variant_size<Shape> : std::integral_constant<std::size_t, 2> {};

    template <>
    struct variant_alternative<Shape::Circle, Shape> { using type = Circle; };

    template <>
    struct variant_alternative<Shape::Rectangle, Shape> { using type = Rectangle; };
}

```

```

Shape::Kind index(const Shape& shape) { return shape.kind; }

template <Kind K>
auto&& get(const Shape& shape) {
    return static_cast<const std::variant_alternative_t<K, Shape>&>(shape);
}

int get_area(const Shape& shape) {
    inspect (shape) {
        <Circle> c: return 3.14 * c.radius * c.radius;
        <Rectangle> r: return r.width * r.height;
    }
    // switch (index(shape)) {
    //     case Shape::Circle: {
    //         const std::variant_alternative_t<Shape::Circle, Shape>& c =
    //             get<Shape::Circle>(shape);
    //         return 3.14 * c.radius * c.radius;
    //     }
    //     case Shape::Rectangle: {
    //         const std::variant_alternative_t<Shape::Rectangle, Shape>& r =
    //             get<Shape::Rectangle>(shape);
    //         return r.width * r.height;
    //     }
    // }
}

```

9.3 Matcher: any_of

The logical-or pattern in other languages is typically spelled *pattern₀ | pattern₁ | ... | pattern_N*, and matches value *v* if any *pattern_i* matches *v*.

This provides a restricted form (constant-only) of the logical-or pattern.

```

template <typename... Ts>
struct any_of : std::tuple<Ts...> {
    using tuple::tuple;

    template <typename U>
    bool match(const U& u) const {
        return std::apply([&](const auto&... xs) { return (... || xs == u); }, *this);
    }
};

```

```

int fib(int n) {
    inspect (n) {
        x if (x < 0): return 0;
        ~(any_of{1, 2}): return n; // 1 | 2
        x: return fib(x - 1) + fib(x - 2);
    }
}

```


9.4 Matcher: within

The range pattern in other languages is typically spelled `first..last`, and matches `v` if `v ∈ [first, last]`.

```
struct within {
    int first, last;

    bool match(int n) const { return first <= n && n <= last; }
};

inspect (n) {
    ^(within{1, 10}): { // 1..10
        std::cout << n << " is in [1, 10].";
    }
    -: {
        std::cout << n << " is not in [1, 10].";
    }
}
```

9.5 Extractor: both

The logical-and pattern in other languages is typically spelled `pattern0 & pattern1 & ... & patternN`, and matches `v` if all of `patterni` matches `v`.

This extractor emulates binary logical-and with a `std::pair` where both elements are references to value `v`.

```
struct Both {
    template <typename U>
    std::pair<U&&, U&&> extract(U&& u) const {
        return {std::forward<U>(u), std::forward<U>(u)};
    }
} both;
```

```
inline constexpr Both both;
```

```
inspect (v) {
    (both! [[x, 0], [0, y]]): // ...
}
```

9.6 Extractor: at

The binding pattern in other languages is typically spelled `identifier @ pattern`, binds `identifier` to `v` and matches if `pattern` matches `v`. This is a special case of the logical-and pattern (`pattern0 & pattern1`) where `pattern0` is an `identifier`. That is, `identifier & pattern` has the same semantics as `identifier @ pattern`, which means we get `at` for free from `both` above.

```
inline constexpr at = both;
```

```
inspect (v) {
    <Point> (at! [p, [x, y]]): // ...
    // ...
}
```

9.7 Red-black Tree Rebalancing

Dereference patterns frequently come into play with complex patterns using recursive variant types. An example of such a problem is the rebalance operation for red-black trees. Using pattern matching this can be expressed succinctly and in a way that is easily verified visually as having the correct algorithm.

Given the following red-black tree definition:

```
enum Color { Red, Black };

template <typename T>
struct Node {
    void balance();

    Color color;
    std::shared_ptr<Node> lhs;
    T value;
    std::shared_ptr<Node> rhs;
};
```

The following is what we can write with pattern matching:

```
template <typename T>
void Node<T>::balance() {
    *this = inspect (*this) {
        // left-left case
        //
        //      (Black) z          (Red) y
        //      /      \          /      \
        //    (Red) y    d      (Black) x (Black) z
        //    /      \          /      \
        //  (Red) x    c      a      b    c      d
        //  /      \
        // a        b
        [^Black, *[^Red, *[^Red, a, x, b], y, c], z, d]
        => Node{Red, std::make_shared<Node>(Black, a, x, b),
                y,
                std::make_shared<Node>(Black, c, z, d)},
        [^Black, *[^Red, a, x, *[^Red, b, y, c]], z, d] // left-right case
        => Node{Red, std::make_shared<Node>(Black, a, x, b),
                y,
                std::make_shared<Node>(Black, c, z, d)},
        [^Black, a, x, *[^Red, *[^Red, b, y, c], z, d]] // right-left case
        => Node{Red, std::make_shared<Node>(Black, a, x, b),
                y,
                std::make_shared<Node>(Black, c, z, d)},
        [^Black, a, x, *[^Red, b, y, *[^Red, c, z, d]]] // right-right case
        => Node{Red, std::make_shared<Node>(Black, a, x, b),
                y,
                std::make_shared<Node>(Black, c, z, d)},
        self => self // do nothing
    };
}
```

The following is what we currently need to write:

```
template <typename T>
void Node<T>::balance() {
    if (color != Black) return;
    if (lhs && lhs->color == Red) {
        if (const auto& lhs_lhs = lhs->lhs; lhs_lhs && lhs_lhs->color == Red) {
            // left-left case
            //
            //      (Black) z          (Red) y
            //      /      \          /      \
            //    (Red) y    d      (Black) x (Black) z
            //   /      \      -> /      \   /      \
            // (Red) x    c      a      b   c      d
            // /      \
            // a      b
            *this = Node{
                Red,
                std::make_shared<Node>(Black, lhs_lhs->lhs, lhs_lhs->value, lhs_lhs->rhs),
                lhs->value,
                std::make_shared<Node>(Black, lhs->rhs, value, rhs)};
            return;
        }
        if (const auto& lhs_rhs = lhs->rhs; lhs_rhs && lhs_rhs->color == Red) {
            *this = Node{ // left-right case
                Red,
                std::make_shared<Node>(Black, lhs->lhs, lhs->value, lhs_rhs->lhs),
                lhs_rhs->value,
                std::make_shared<Node>(Black, lhs_rhs->rhs, value, rhs)};
            return;
        }
    }
    if (rhs && rhs->color == Red) {
        if (const auto& rhs_lhs = rhs->lhs; rhs_lhs && rhs_lhs->color == Red) {
            *this = Node{ // right-left case
                Red,
                std::make_shared<Node>(Black, lhs, value, rhs_lhs->lhs),
                rhs_lhs->value,
                std::make_shared<Node>(Black, rhs_lhs->rhs, rhs->value, rhs->rhs)};
            return;
        }
        if (const auto& rhs_rhs = rhs->rhs; rhs_rhs && rhs_rhs->color == Red) {
            *this = Node{ // right-right case
                Red,
                std::make_shared<Node>(Black, lhs, value, rhs->lhs),
                rhs->value,
                std::make_shared<Node>(Black, rhs_rhs->lhs, rhs_rhs->value, rhs_rhs->rhs)};
            return;
        }
    }
}
```

10 Future Work

10.1 Language Support for Variant

The design of this proposal also accounts for a potential language support for variant. It achieves this by keeping the alternative pattern flexible for new extensions via `< new_entity > pattern`.

Consider an extension to `union` that allows it to be tagged by an integral, and has proper lifetime management such that the active alternative need not be destroyed manually.

```
// `: type` specifies the type of the underlying tag value.  
union U : int { char small[32]; std::vector<char> big; };
```

We could then allow `< qualified-id >` that refers to a `union` alternative to support pattern matching.

```
U u = /* ... */;  
  
inspect (u) {  
    <U::small> s: std::cout << s;  
    <U::big> b: std::cout << b;  
}
```

The main point is that whatever entity is introduced as the discriminator, the presented form of alternative pattern should be extendable to support it.

10.2 Note on Ranges

The benefit of pattern matching for ranges is unclear. While it's possible to come up with a ranges pattern, e.g., `{x, y, z}` to match against a fixed-size range, it's not clear whether there is a worthwhile benefit.

The typical pattern found in functional languages of matching a range on head and tail doesn't seem to be all that common or useful in C++ since ranges are generally handled via loops rather than recursion.

Ranges likely will be best served by the range adaptors / algorithms, but further investigation is needed.

11 Acknowledgements

Thanks to all of the following:

- Yuriy Solodkyy, Gabriel Dos Reis, Bjarne Stroustrup for their prior work on [N3449], Open Pattern Matching for C++ [OpenPM], and the [Mach7] library.
- Pattern matching presentation by Bjarne Stroustrup at Urbana-Champaign 2014. [PatMatPres]
- Jeffrey Yasskin/JF Bastien for their work on [P1110R0].
- (In alphabetical order by last name) Dave Abrahams, John Bandela, Agustín Bergé, Ori Bernstein, Matt Calabrese, Alexander Chow, Louis Dionne, Michał Dominiak, Vicente Botet Escribá, Eric Fiselier, Bengt Gustafsson, Zach Laine, Jason Lucas, John Skaller, Bjarne Stroustrup, Tony Van Eerd, and everyone else who contributed to the discussions.

12 References

- [Mach7] Yuriy Solodkyy, Gabriel Dos Reis, and Bjarne Stroustrup. Mach7: Pattern Matching for C++. <https://github.com/solodon4/Mach7>
- [N3449] Bjarne Stroustrup. 2012. Open and Efficient Type Switch for C++. <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2012/n3449.pdf>
- [OpenPM] Yuriy Solodkyy, Gabriel Dos Reis, and Bjarne Stroustrup. Open Pattern Matching for C++. <http://www.stroustrup.com/OpenPatternMatching.pdf>
- [P0095R0] David Sankel. 2015. Pattern Matching and Language Variants. <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2015/p0095r0.html>
- [P0144R2] Herb Sutter, Bjarne Stroustrup, and Gabriel Dos Reis. 2016. Structured bindings. <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2016/p0144r2.pdf>
- [P1110R0] Jeffrey Yasskin and JF Bastien. 2018. A placeholder with no name. <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2018/p1110r0.html>
- [P1260R0] Michael Park. 2018. Pattern Matching. <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2018/p1260r0.pdf>
- [P1308R0] David Sankel, Dan Sarginson, and Sergei Murzin. 2018. Pattern Matching. <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2018/p1308r0.html>
- [PatMatPres] Yuriy Solodkyy, Gabriel Dos Reis, and Bjarne Stroustrup. “Pattern Matching for C++” presentation at Urbana-Champaign 2014.
- [Patterns] Michael Park. Pattern Matching in C++. <https://github.com/mpark/patterns>
- [SimpleMatch] John Bandela. Simple, Extensible C++ Pattern Matching Library. https://github.com/jbandela/simple_match
- [Swift Patterns] Swift Reference Manual - Patterns. <https://docs.swift.org/swift-book/ReferenceManual/Patterns.html>
- [Warnings] Luc Maranget. Warnings for pattern matching. <http://moscova.inria.fr/~maranget/papers/warn/index.html>
- [YAMLParser] http://llvm.org/doxygen/YAMLParser_8h_source.html